

Analysis of Adobe Flash Player ID3 Tag Parsing Integer Overflow Vulnerability (CVE-2015-5560)□

January 12, 2016 by Nahuel Riva

Vulnerability Overview

After Adobe released a [patch for this vulnerability](#), it was made public that this bug [was already being exploited](#) in the wild by some exploit kits like *Angler* and *Nuclear Pack*.

This vulnerability is about an integer overflow in *Adobe Flash Player* when parsing a compressed [ID3 tag](#) which size exceed `0x2AAAAAAAA` bytes. An error in how the size of a dynamic allocated buffer is calculated, used as destination for final decompressed data, produces that too much data is copied to a small buffer. In other words, a heap-based buffer overflow.

This bug was fixed in *Adobe Flash player 18.0.0.232*. However, this is an important fix because that version, the previous one *18.0.0.209*, and new versions, [introduce new exploit mitigations](#) to avoid exploitation techniques as the one described in [Haife's Li presentation](#) using *Vector*.

In fact, the exploits included in the exploit kits mentioned above, perform new bypasses as the ones described in the [Project Zero blog spot](#).

Vulnerability analysis, finding the root cause□

[Natalie Silvanovich](#), from *Google Project Zero (PZ)*, [made public a PoC](#) in order to trigger this bug. As the PoC seems to be written for Flash CS (or some compliant compiler) and I like more Apache Flex, I re-wrote the PoC like this:

```

package
{
    import flash.display.*;
    import flash.media.*;
    import flash.utils.*;
    import flash.net.*;
    import flash.events.*;
    import flash.system.*;
    import flash.external.*;
    import avm2.intrinsics.memory.*;

    public class CVE_2015_5560 extends Sprite
    {
        public var mySound:Sound;

        function CVE_2015_5560()
        {
            logDebug("Loading MP3 file ...");

            mySound = new Sound();
            mySound.load(new URLRequest("CVE_2015_5560.mp3"));
            mySound.play();

            logDebug("Triggering corruption :)");

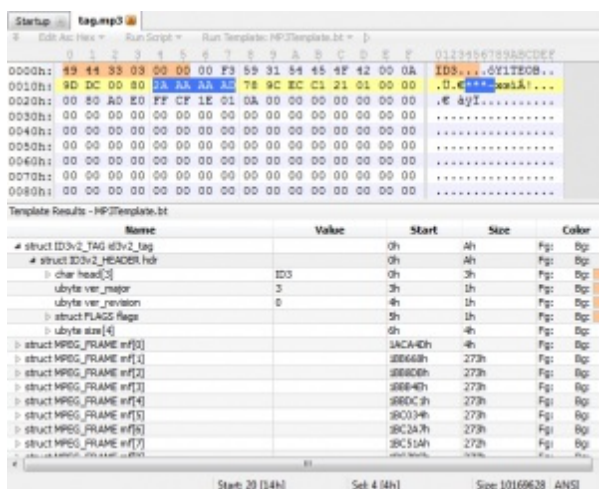
            setInterval(f, 1000);
        }

        private function f():void
        {
            mySound.id3;
        }
    }
}

```

Anyway, the important thing here is the MP3 referenced in the code; that's the real file triggering the bug.

Let's see a little bit the structure of the ID3 tag contained in the MP3 file with the help of the 010 Editor and its MP3 template:



We can observe the presence of an ID3v2 tag followed by another tag that the template couldn't recognize due to malformed data:

```
MP3: ID3v2 tag found
MP3: warning: invalid ID3v2 tag header --> ERROR HERE!!!
MP3: warning: invalid MPEG frame synchronization at offset 0xA
MP3: warning: invalid MPEG header in frame at offset 0x1ACA4D
MP3: warning: invalid MPEG frame synchronization at offset 0x1ACA51
MP3: first found MPEG frame parameters:
MP3: - header offset: 0x1BB668
MP3: - bitrate: 192 kbit
MP3: - MPEG-1 layer 3
MP3: - sampling frequency: 44100 Hz
MP3: - channel mode: stereo
MP3: - CRC protected: No
MP3: ID3v1 tag found
MP3: file parsing completed!
MP3: valid MPEG frames found: 13324
MP3: average frame bitrate: 192 kbit
```

In the previous image, we can see the **major** and **revision** fields in the *ID3* header with the values **3** and **0**, respectively, which means that this is an ID3v2.3.0. Hence I'm going to use that version of the [specs](#) to perform my analysis of the tag.

According to the specification, after the main **ID3 header**, we can find the so called **ID3 frames**, with the following structure:

```
Frame ID  $xx xx xx xx (four characters)
Size      $xx xx xx xx
Flags     $xx xx
```

So, we have 10 bytes, adjusting this to what we have seen previously, our tag ends like this:

Frame ID: 'TEOB'
Size: 0x000A9DDC
Flags: 0x0080

First, we have the **Frame ID**. In our case, it is a [user defined text information frame](#).

After the **Frame ID**, we have the **Size**, *0xA9DDC*. This is the *frame size* excluding the *frame header*, that is *frame size* – 10.

At the end, we have the **Flags**, **0x0080**. This field is composed by two bytes, **0x00** and **0x80**. The first byte, *0x00*, is the **status message** and the second one, *0x80*, is used for **encoding purposes**.

According to the mentioned in the [3.3.1. Frame header flags](#) section, our first byte is *0x00*, so, our result is:

- **Tag alter preservation**: Frame should be preserved.
- **File alter preservation**: Frame should be preserved.
- **Read only**: no read-only.

These are not important flags for us, let's see the other byte, *0x80* (*10000000* in binary):

- **Compression**: 1 Frame is compressed using [zlib](#) with 4 bytes for 'decompressed size' appended to the frame header.
- **Encryption**: 0 Frame is not encrypted.
- **Grouping identity**: 0 Frame does not contain group information.

In this second byte, only the most significant bit, the **Compression** bit, is on so that indicates that our frame is compressed using **zlib**. Therefore, according to the specs, there must be **decompression size** (4 bytes) field after the *frame header*. In our case, our *decompression size* is *0x2AAAAAAA*. What's next is the compressed data using *zlib*. Just to be sure, I wrote a small and ugly script to decompress the data:

```

import sys
import zlib

from struct import unpack

print 'Opening file ...'
fd = open(sys.argv[1], 'rb')
data = fd.read()
fd.close()

print 'Reading tag size ...'
tag_size = unpack('>L', data[0x0E:0x0E+4])[0]
print 'Tag size: %x' % tag_size

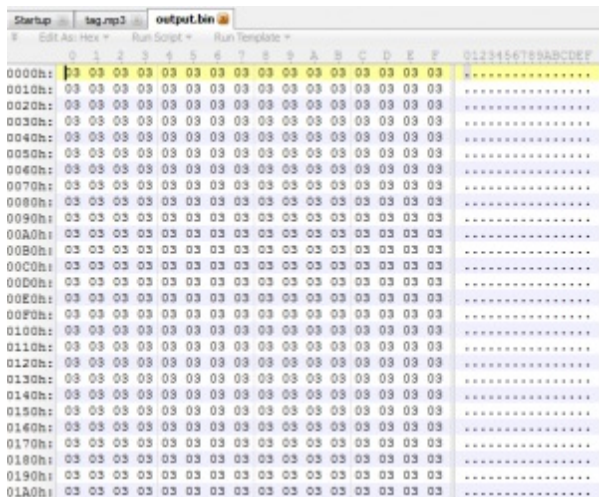
print 'Getting compressed data ...'
compressed_data = data[0x18:0x18+tag_size]

print 'Decompressing ...'
decompressed_data = zlib.decompress(compressed_data)

print 'Saving decompressed data to a file ...'
fd = open('output.bin', 'wb')
fd.write(decompressed_data)
fd.close()
print 'Done. Saved to output.bin'

```

This is the result:



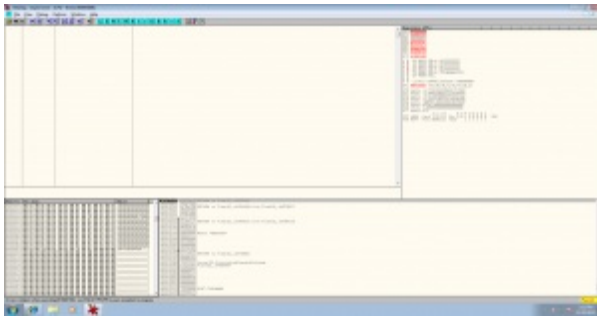
The decompressed data contains a lot of 0x03 bytes.

According to the *PZ* advisory, the *decompression size* is the value that causes the integer overflow so, it's important to pay a attention to this value during our debugging session.

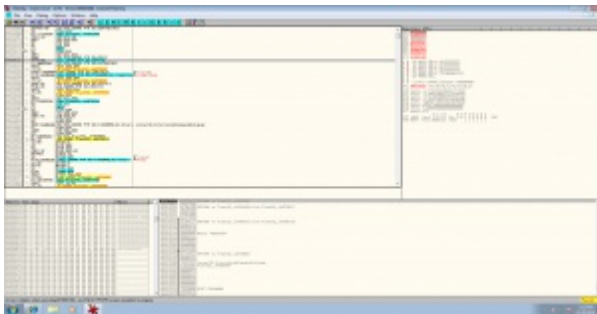
Just to see how *Adobe Flash Player* behaves using this *PoC*, I compiled the *AS* code and put the generated *SWF* file plus the *MP3* file and an *HTML* file that loads the *SWF* in a folder and then started a webserver using *Python* like this: `python -m SimpleHTTPServer 8888`.

My testing environment is Windows Ultimate SP1 with *Adobe Flash Player* 18.0.0.209 (32 bits).

Once I requested the *HTML* and the *SWF* was loaded, this was the result:



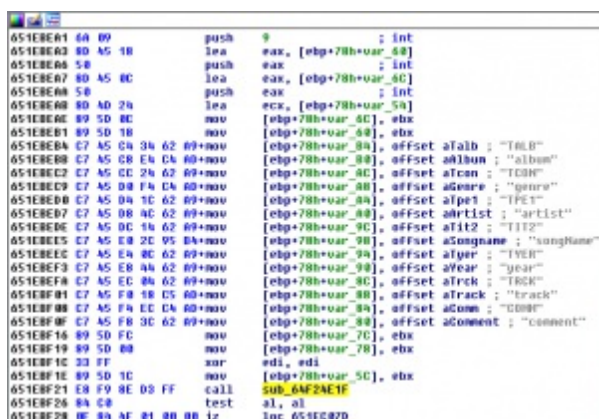
In the previous image, *EIP* has the value *0x01000100* and in *EAX* we can recognize a known value, *0x03030303* (compressed data). In this case, we were a lucky guys because *0x03030303* is a mapped address and its content ended up in *EIP*. When this address is not mapped, the program crashes earlier, in a *CALL [EAX+8]*, just in the virtual call of the *id3* property from the **mySound** object.



Next step is to identify where the integer overflow is produced so we must find the function responsible for parsing the *ID3* tag.

Generally, an *ID3* tag is used to show meta-information as the album name, music genre, artist name, etc; of a *MP3* file (or any audiovisual container file) so I started to look for strings in *IDA* like “album“, “genre“, “author“, etc.

Through the “genre” string, I finally found a nice basic block where the tag seems to be parsed:



The *sub_64F24E1F* function is responsible for parsing all the frames in the archive. We can set

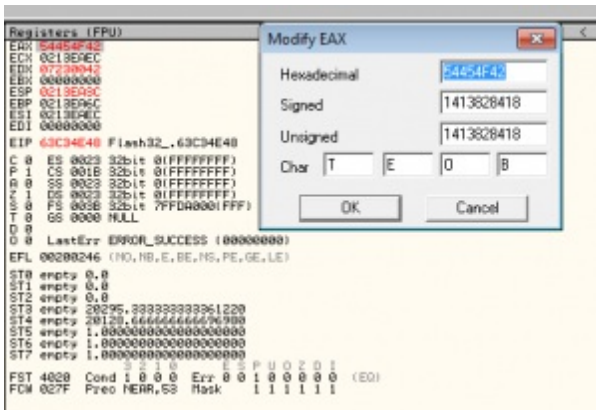
a breakpoint at the very beginning of the function and start tracing to identify the important pieces of the code.

Here's a brief summary of the most important pieces of the code. Just for convenience, I renamed the function to `parse_id3_tag`.

What happens first in `parse_id3_tag+24`, when calling `GetFrameInfo`, is that the *Frame ID* is read:

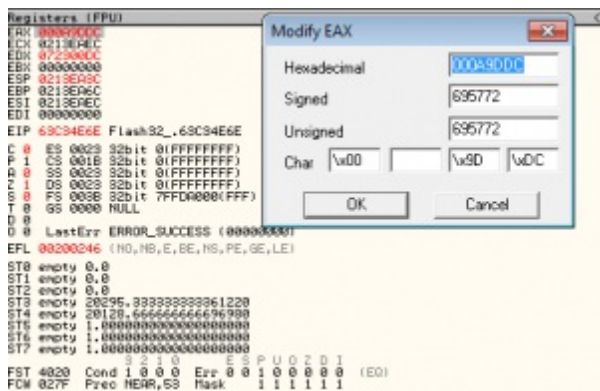
```
parse_id3_tag+1D  loc_64F24E3C:
parse_id3_tag+1D  xor     ebx, ebx
parse_id3_tag+1F  push   ebx
parse_id3_tag+20  push   4
parse_id3_tag+22  mov     ecx, esi
parse_id3_tag+24  call   GetFrameInfo
parse_id3_tag+29  mov     [ebp+var_14], eax
parse_id3_tag+2C  cmp     eax, ebx
parse_id3_tag+2E  jz      loc_64F25160
```

In EAX, we have the return value:



After that, the *Size* is read:

```
parse_id3_tag+43
parse_id3_tag+43  loc_64F24E62:
parse_id3_tag+43  movzx   eax, byte ptr [esi+20h]
parse_id3_tag+47  push   eax
parse_id3_tag+48  push   4
parse_id3_tag+4A  call   GetFrameInfo
```



Then, we get the *Flags*:

```

parse_id3_tag+61  push    ebx
parse_id3_tag+62  push    1
parse_id3_tag+64  mov     ecx, esi
parse_id3_tag+66  call    GetFrameInfo >> encoding
parse_id3_tag+6B  push    ebx
parse_id3_tag+6C  mov     edi, eax
parse_id3_tag+6E  push    1
parse_id3_tag+70  shl     edi, 8
parse_id3_tag+73  call    GetFrameInfo >> status message
parse_id3_tag+78  mov     ecx, [ebp+var_8]
parse_id3_tag+7B  mov     ebx, eax
parse_id3_tag+7D  mov     al, [esi+20h]
parse_id3_tag+80  or      ebx, edi
parse_id3_tag+82  cmp     byte ptr [esi+21h], 4
parse_id3_tag+86  mov     [ebp+var_20], ebx
parse_id3_tag+89  mov     [ebp+var_10], al
parse_id3_tag+8C  mov     [ebp+var_1C], ecx
parse_id3_tag+8F  mov     [ebp+var_1], 1
parse_id3_tag+93  jnz     short loc_64F24F12

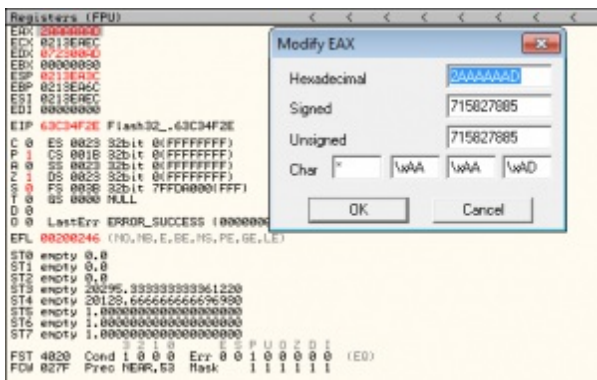
```

As we mentioned earlier, the *Compression* bit is on so there must be a call to get the *decompression size* value and that's exactly what happens next:

```

parse_id3_tag+103  push    dword ptr [ebp+var_10]
parse_id3_tag+106  mov     ecx, esi
parse_id3_tag+108  push    4
parse_id3_tag+10A  call    GetFrameInfo
parse_id3_tag+10F  sub     [ebp+var_8], 4
parse_id3_tag+113  mov     [ebp+var_1C], eax

```

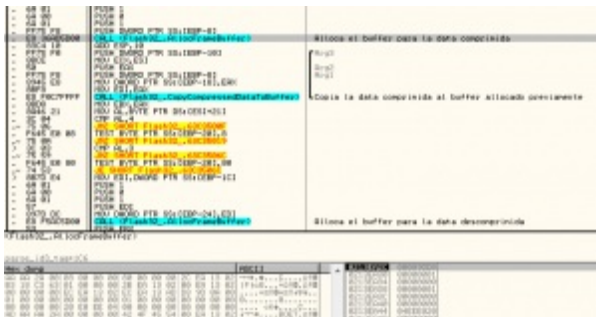
Now that the frame size and the size for the decompressed data have been read, it's time to allocate the corresponding buffers to hold the data.

Starting at `parse_id3_tag+1BD`, it allocates a buffer using `frame_size - 4` as the size for the buffer:

```

parse_id3_tag+1BD  push  1          ; char
parse_id3_tag+1BF  push  0          ; int
parse_id3_tag+1C1  push  1          ; int
parse_id3_tag+1C3  push  [ebp+var_8]    ; int >> frame size
parse_id3_tag+1C6  call  AllocFrameBuffer
parse_id3_tag+1CB  add    esp, 10h
parse_id3_tag+1CE  push  dword ptr [ebp+var_10] ; char
parse_id3_tag+1D1  mov    ecx, esi
parse_id3_tag+1D3  push  eax          ; int >> buffer to store data
parse_id3_tag+1D4  push  [ebp+var_8]    ; int >> frame size - 4
parse_id3_tag+1D7  mov    [ebp+var_18], eax
parse_id3_tag+1DA  mov    edi, eax
parse_id3_tag+1DC  call  CopyCompressedDataToBuffer
parse_id3_tag+1E1  mov    ebx, eax
parse_id3_tag+1E3  mov    al, [esi+21h]
parse_id3_tag+1E6  cmp    al, 4
parse_id3_tag+1E8  jnz    short loc_64F2500F

```



Remember that the size was `0x0A9DDC` and now it is `0x0A9DD8`.

In this case, `0x07A39000` is the buffer used to hold the compressed data. This data is copied there using the `CopyCompressedDataToBuffer` function:

[illegible]

After that, it allocates a buffer using the *decompression size* and then decompress and copy the data to the buffer:

```

parse_id3_tag+1FA      parse_id3_tag+1FA      loc_64F25019:
parse_id3_tag+1FA      mov     edi, [ebp+var_1C]
parse_id3_tag+1FD      push    1             ; char
parse_id3_tag+1FF      push    0             ; int
parse_id3_tag+201      push    1             ; int
parse_id3_tag+203      push    edi           ; int >> decompression size
parse_id3_tag+204      mov     [ebp+var_24], edi
parse_id3_tag+207      call   AllocFrameBuffer
parse_id3_tag+20C      push    ebx           ; int >> frame size - 4
parse_id3_tag+20D      push    [ebp+var_18]   ; int >> buffer with compressed data
parse_id3_tag+210      lea     ecx, [ebp+var_24]
parse_id3_tag+213      push    ecx           ; int
parse_id3_tag+214      push    eax           ; int >> buffer to store the data
parse_id3_tag+215      mov     [ebp+var_20], eax
parse_id3_tag+218      call   DecompressZlibData
parse_id3_tag+21D      add     esp, 20h
parse_id3_tag+220      test    eax, eax
parse_id3_tag+222      jnz     loc_64F25160

```

[illegible]

At this point, we don't notice anything unusual, everything seems to be fine. We keep in mind that in `0x07C30000` we have the decompressed data but a little bit further, in `parse_id3_tag+281`, we see this:

```

parse_id3_tag+281
parse_id3_tag+281  loc_64F250A0:
parse_id3_tag+281  mov     eax, ebx
parse_id3_tag+283  imul    eax, 6
parse_id3_tag+286  add     eax, 2
parse_id3_tag+289  cmp     [esi+28h], eax
parse_id3_tag+28C  mov     [ebp+var_20], eax
parse_id3_tag+28F  jge     short loc_64F250DA

```

In *EBX*, we have the *decompression size* – 1 (this has to do with a little code that we overlooked [1]), *0x2AAAAAAC*. Then, *EBX* is copied to *EAX* and multiplied by 6. The result is stored in *EAX*.

Let's do some math:

$$(0x2AAAAAAC * 6) + 2 = 0x10000000A$$

Now we do find the bug, the integer overflow is very obvious, what only fits in 32 bits is *0x0A*. Then, the result of this operation is used as a size to allocate a buffer in *parse_id3_tag+2A6*:

```

parse_id3_tag+2A6
parse_id3_tag+2A6  loc_64F250C5:      ; char
parse_id3_tag+2A6  push    1
parse_id3_tag+2A8  push    0      ; int
parse_id3_tag+2AA  push    1      ; int
parse_id3_tag+2AC  push    eax      ; int >> overflowed size
parse_id3_tag+2AD  mov     [esi+28h], eax
parse_id3_tag+2B0  call   AllocFrameBuffer
parse_id3_tag+2B5  add     esp, 10h
parse_id3_tag+2B8  mov     [esi+24h], eax

```

In this case, the buffer starts at *0x04A7C920* and its size is, as we shown before, *0x0A*.

Then, in *parse_id3_tag+2D5*, there is a function call to copy the remaining decompressed data to the previously allocated buffer with the overflowed size:

```

parse_id3_tag+2CA  push    [ebp+var_1C] ; int
parse_id3_tag+2CD  lea     eax, [ebx+edi]
parse_id3_tag+2D0  push    eax      ; int >> value used as MAX counter(EBX+EDI)
parse_id3_tag+2D1  push    edi      ; int >> src buffer (with decompressed data)
parse_id3_tag+2D2  push    ecx      ; int >> dest buffer (allocated with overflowed size)
parse_id3_tag+2D3  mov     ecx, esi
parse_id3_tag+2D5  call   InternalMemcpy
parse_id3_tag+2DA  mov     ecx, [ebp+var_20]

```

If we look it in Ollydbg, we see this:

0213EA2C	04A7C920	Arg1 = 04A7C920
0213EA30	07C30001	Arg2 = 07C30001
0213EA34	326DAAAD	Arg3 = 326DAAAD
0213EA38	00000003	Arg4 = 00000003
0213EA3C	00000000	
0213EA40	00000000	
0213EA44	04EEE020	
0213EA48	2AAAAAAD	
0213EA4C	04A7C920	
0213EA50	00000003	
0213EA54	07C30000	

We can see that one of the parameters used in the function call is `0x326DAAAD`. This is the result of executing the `lea eax, [ebx+edi]` instruction located at `parse_id3_tag+2CD`. *EBX* had *decompression size - 1* while *EDI* pointed to the buffer where the *decompressed data + 1* was, so, `0x02AAAAAAC + 0x07C30001 = 0x326DAAAD`. This value it's going to be used in `InternalMemcpy+D8` as one of the conditions in the loop. The other condition in the loop is to copy until a null byte is found.

After tracing some rounds of the loop, we can see how the copy operation was performed beyond the `0x0A` bytes:

Address	Hex dump	ASCII
04A7C921	03 03 03 03 03 03 03 03 03 03 03 03 03 03 03 03	*****
04A7C925	03 03 03 03 03 03 03 03 03 03 03 03 03 03 03 03	*****
04A7C929	03 03 03 03 03 03 03 03 03 03 03 03 03 03 03 03	*****
04A7C92D	03 03 03 03 03 03 03 03 03 03 03 03 03 03 03 03	*****
04A7C931	03 03 03 03 03 03 03 03 03 03 03 03 03 03 03 03	*****
04A7C935	03 03 03 03 03 03 03 03 03 03 03 03 03 03 03 03	*****
04A7C939	03 03 03 03 03 03 03 03 03 03 03 03 03 03 03 03	*****
04A7C93D	03 03 03 03 03 03 03 03 03 03 03 03 03 03 03 03	*****
04A7C941	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	*****
04A7C945	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	*****
04A7C949	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	*****
04A7C94D	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	*****
04A7C951	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	*****
04A7C955	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	*****


Boom, heap overflow in sight!

Some words about the exploitation process


To summarize, the root vulnerability is an integer overflow that then ends up in a heap-based buffer overflow. In a very basic exploitation scenario, we must massage the heap in order to place a buffer with our data just after the buffer we overflow. Also, as we have protections such as *ALSR/DEP/CFG* that we must bypass in order to get reliable code execution, it would be a good idea to build some kind of read/write primitives that help us during the exploitation process.

For *Flash 18.0.0.209*, the most used technique to exploit *Adobe Flash* was the one described by **Haifei Li**, entitled “[Smashing the Heap with Vector: Advanced Exploitation Technique in Recent Flash Zero-day Attack](#)“. By incrementing the value stored in the *length* field of a *Vector.<uint>* object, we can read and write beyond the limits of the *Vector*.

So, we could place a *Vector.<uint>* after our overflowed buffer and overwrite the metadata, the *length* field to be more accurate, and build our read/write primitive. The following picture (borrowed from *Project's Zero* post):



During the exploitation process, it ends up like this:



But starting from *Adobe Flash 18.0.0.209* we have some new challenges to face before getting reliable code execution because new exploit mitigation mechanisms were added in order to avoid the mentioned method from *Haifei Li*. For example, now *Flash* has some heap isolation mechanism called **heap partitioning** in which objects like *Vector.<uint>* are isolated from the *Flash* heap and stored in the **System heap**. So, our exploitation scenario turns into something like this:



As if this were not enough, now we have better ASLR in the heap. To highlight:

- Allocations > 1 MB have better randomization
- In x64, the *Flash* heap ends up far away of any mapped area

Last but not least, a validation to the length fields of the *Vector.<*>* object was added.

If you want to read more detailed information about all these new mechanisms introduced in *Flash* you can go to the [excellent write up](#) published by *Project Zero*.

So, the only thing to think at this point is: we are screwed!.

In order to exploit this bug in a reliable way, at least in *Flash 18.0.0.209*, we are going to not only allocate in a deterministic way (to avoid randomization) an object that places after our overflowed buffer (*Flash* heap) but also it must have the capability to give us the possibility (as the length of a *Vector.<uint>* object did it before) to generate a read/write primitive.

So, please, stay tuned to see how this story ends!

Notes

- [1] We only get to the zone where the overflow occurs just if the *JG* jump located at *parse_id3_tag+27B* is not executed, that is, if the condition is *False*. Once the data was decompressed it takes the first byte of the buffer, it decrements the *decompression size* by 1 and compares it with 3. If it is greater, it goes to the end of the loop and process the next frame. If not, it goes to the overflow zone.

 Share < 110

 Tweet < 291

 Email < 3

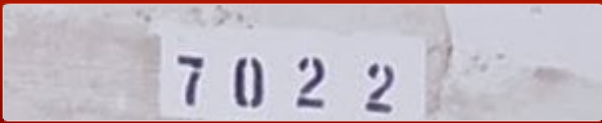
 Share < 757

 G+1 < 1




Latest from CoreLabs


- ◀ Core Security Selected as SC Magazine 2016 Excellence Award Finalist

Leave a Comment



Privacy & Terms





Post Comment

[Blog Home](#)

[Latest from CoreLabs](#)

[What's New At Core](#)

Resources

[Case Studies](#)

[Data Sheets](#)

[Quick Reads](#)

[Webcasts](#)

[White Papers](#)

[Videos](#)

Core Impact Pro Demo

CORE IMPACT PRO®

[Request A Demo Now](#)

Core News Sign Up

Name

First

Last

Email *

Agreement

☐ I agree to receive information regarding Core products and services.

[Submit](#)

RSS Feed

Subscribe to our [RSS feed](#) to get the latest blog posts - formatted for your favorite feed reader



© Core Security 2016